

UNIVERSITY OF PISA



MULTIMEDIA INFORMATION RETRIEVAL AND
COMPUTER VISION

PROJECT

Search Engine

Bergami Giovanni

Bruchi Caterina

Grillea Francesco

A.Y. 2023-2024

github.com/francescogrillea/SearchEngine_MIRCVProject

Contents

1	Overview	2
1.1	Data Folder	2
1.2	Project Structure	3
2	Data Structures	4
2.1	Posting List	4
2.1.1	Skipping Pointer	5
2.2	Lexicon	6
2.2.1	TermEntryList	6
2.2.2	TermEntry	7
2.3	Document Index	8
2.3.1	DocInfo	8
2.4	Readers and Writers	8
2.4.1	PostingListReader	8
2.4.2	PostingListBlockReader	8
2.4.3	LexiconReader	10
2.4.4	DocIndexReader	10
2.5	SPIMI Algorithm	10
3	Experimental Results	11
3.1	In-Memory Structures	12

Chapter 1

Overview

1.1 Data Folder

```
data/
├── analysis/
│   ├── memoryDump-MainClass.hprof
│   └── memoryDump-MainClass[-p].hprof
├── evaluation/
│   ├── trec_eval-9.0.7
│   ├── 2020qrels-pass.txt
│   └── msmarco-test2020-queries.tsv.gz
├── intermediate_postings/
│   ├── doc_index/
│   ├── index/
│   └── lexicon/
├── lexicon.bin
├── doc_index.bin
├── index.bin
├── cleanup.sh
├── collection.tar.gz
├── collection_subset.sh
└── stop_words_english.txt
```

where

- `analysis` folder contains two memory dump of the query execution phase in order to prove that, even keeping the whole lexicon in main memory, the occupation of the whole main class is below 512MB (see Chapter 3.1).
- `evaluation` folder are stored all the components to produce the results

given by the evaluation phase.

- `intermediate_postings` contains the whole intermediate files generated during the SPIMI algorithm.
- `lexicon.bin`, `doc_index.bin` and `index.bin` are the final data structures stored in a suitable format generated in the merge phase of the SPIMI algorithm. Are the files used by the query execution phase.
- `cleanup.sh` is an utility bash script that removes all data structures stored on disk (even the intermediate ones).
- `collection.tar.gz` the dataset.
- `collection_subset.sh` is an utility bash script that generates a subset of `n` elements of the dataset. Adding the flag `-r`, the documents will be picked randomly.
- `stop_words_english.txt` is the list of 851 english stopwords used in stopword removal process.

1.2 Project Structure

```
src
├── main
│   ├── java/org
│   │   ├── common
│   │   │   └── encoding
│   │   ├── evaluation
│   │   ├── offline_phase
│   │   ├── online_phase
│   │   │   ├── query_processing
│   │   │   └── scoring
│   └── test
│       └── java/org
```

Chapter 2

Data Structures

In the following section we're going to take an overview of all data structures used in this project.

2.1 Posting List

The `PostingList` class represents a posting list for a given term in lexicon and is made up by the following variables and methods.

Instance Variables

- `List<Integer> doc_ids`: a list of document IDs associated with the terms in the posting list.
- `List<Integer> term_frequencies`: a list of term frequencies, representing the number of occurrences of each term in the corresponding documents.
- `List<SkippingPointer> skipping_points`: a list of skipping pointers, which are used to optimize search operations by skipping over a certain number of entries.
- `int block_size`: indicating the number of elements between two skipping pointers. Is assigned in the `initPointers()` method to the squared root of the number of postings.

Methods

In addition to getter, setters, `toString`, we defined the following additional methods:

- `void addPosting(int doc_id)`: add `doc_id` to the posting list increasing its term frequency if the `doc_id` is already present.
- `void appendPostings(PostingList new_postings)`: appends the `new_postings` to the current posting list increasing term frequency of documents that are already present.
- `ByteBuffer serialize()`: is used to serialize the intermediate index, i.e. without skipping pointers and compression. It serializes the docIDs and term frequencies one by one separating the two list by the -1 value.
- `ByteBuffer serializeDocIDsBlock(EncoderInterface encoder, int start_index, int end_index)`: serializes a whole block defined by the subset going from `start_index` to `end_index` of docIDs using the encoder technique passed as input.
- `ByteBuffer serializeTFsBlock(EncoderInterface encoder, int start_index, int end_index)`: serializes a whole block defined by the subset going from `start_index` to `end_index` of term frequencies using the encoder technique passed as input.
- `void initPointers()`: divides the posting list into blocks of size \sqrt{len} .

2.1.1 Skipping Pointer

The `SkippingPointer` class represent the information of a block used for efficient transversal of posting lists. Is made up by the following variables and methods.

Instance Variables

- `max_doc_id`: the maximum docID in the corresponding block.
- `block_length_docIDs`: block size (in bytes) of the corresponding docID list.
- `block_length_TFs`: block size (in bytes) of the corresponding term frequency list.

Methods

- `ByteBuffer serialize()`: is used to serialize the skipping pointer into a `ByteBuffer` in order to write it efficiently on disk.

2.2 Lexicon

The `Lexicon` class implements a mapping between terms and the corresponding posting list's position in index file(s). Is made up by the following variables and methods.

Instance Variables

- `HashMap<String, TermEntryList> lexicon`: the mapping between a term and its corresponding posting list information. Note that for a term we can have multiple `TermEntry` instances: this speeds up computation while merging the intermediate postings lists because we have the location of all posting lists of that term along all the intermediate index files.
- `int size = 0`: number of terms in the lexicon, this is useful in the `add` method in order to keep an association between term and posting list while creating intermediate index.

Methods

- `int add(String term)`: adds a term to the lexicon and returns the index of the corresponding posting list inside the intermediate index (which is a list of posting lists).
- `int add(String term, TermEntryList entries)`: appends a list of term entry to the `TermEntryList` relative to the given term. This is useful for speed up computation while merging the intermediate postings lists.
- `ByteBuffer serializeEntry(String key)`: is used to serialize an single `Lexicon` entry into a `ByteBuffer`.
- `void merge(Lexicon new_lexicon)`: merges the `new_lexicon` with the current one.

2.2.1 TermEntryList

The `TermEntryList` class represents a list of `TermEntry` object. However, in the final merged lexicon, the `TermEntryList` has only one `TermEntry` for each term.

Instance Variables

- `int term_index`: an index associated to the respective term.

- `List<TermEntry> termEntryList`: the list of `TermEntry` instances representing the location of posting list(s) across all the (intermediate and not) index files.

Methods

- `void resetTermEntry(TermEntry termEntry)`: delete all values instances in the `termEntryList` and adds only the one passed as argument.
- `ByteBuffer serialize()`: serialized the first (and only one) `TermEntry` instance inside the list. Note that when the `TermEntryList` contains more than one term entry during the merge phase, it is kept in main memory: writing multiple term entries on disk never happens!

2.2.2 TermEntry

The `TermEntry` class represents metadata about terms in information retrieval systems, such as the position of the posting list of the relative term inside the index file. Is made up by the following instance variables, however besides to getters, setters and `serialize()`, no additional methods are required.

Instance Variables

- `int chunk_index`: the chunk index where the term is located. Useful in merging intermediate postings lists.
- `long offset`: the offset within the file where the term's posting list starts.
- `long length`: the length of the term's postings within the file.
- `int document_frequency`: the document frequency of the term in the collection, i.e., how many documents that term contains.
- `float tfidf_upper_bound`: the upper bound for TF-IDF scoring of the term.
- `float bm25_upper_bound`: the upper bound for BM25 scoring of the term.
- `int BYTES`: the number of bytes required to serialize a `TermEntry` object.

2.3 Document Index

The `DocIndex` class represents the document index structure for storing document information. Is made up by `HashMap<Integer, DocInfo> docIndex` which keeps the association between `doc_id` and its document information; however besides to add, getters, setters and `serialize()`, no additional methods are required. Is important to underline that this data structure is used only while generating intermediate index, in all other cases the document information are retrieved directly by file through the `DocIndexReader` class.

2.3.1 DocInfo

The `DocInfo` class represents information associated with a document. Is made up by the following instance variables, however besides to getters and `serialize()`, no additional methods are required.

- `int pid`: the original identifier of the document.
- `int length`: the length of the (processed) document.
- `int BYTES`: the number of bytes required to represent a serialized `DocInfo` object

2.4 Readers and Writers

For each of the above data structure, ad hoc classes have been created in order to write and read efficiently binary objects from disk.

2.4.1 PostingListReader

The `PostingListReader` class provides utility methods for reading and writing `PostingList` objects to and from an index file specifying the encoding techniques to be used during compression process. Is important to notice that `read` methods take as argument a `TermEntry` in order to locate and read the posting list data directly from the index file; on the other hand `write` methods returns a `TermEntry` instance in which is contained all information of where that posting list has been stored.

2.4.2 PostingListBlockReader

The `PostingListBlockReader` class provides functionality for block-wise reading of posting lists, where each block is accompanied by a `SkippingPointer`.

The reader can iterate through the posting list blocks and retrieve information about document IDs and term frequencies for a given term. Each term has its own `PostingListBlockReader` instance. Is made up by the following instance variables and methods:

Instance Variables

- `String term`: the term itself.
- `FileChannel fileChannel`: the file channel associated to that posting list.
- `PostingList current_block`: the current block read.
- `long last_byte`: the position of the last byte in the current posting list, used to check if the posting list is terminated. Otherwise, we could read postings of another term.
- `int index_pointer`: the index indicating the current position within the current block.
- `float termUpperBound`: the upper bound value associated with the term (e.g., TF-IDF or BM25 upper bound).
- `int documentFrequency`: the document frequency of the term in the entire document collection.

Methods

- `SkippingPointer readBlockPointer()`: reads the skipping pointer associated with the current block.
- `void readBlockContent(SkipingPointer pointer)`: reads the content (docIDs and termFreqs) of the current block.
- `boolean readBlock()`: reads the next block of the posting list. Returns true if a new block is successfully read, false if the end of the posting list is reached. It calls `readBlockPointer` and `readBlockContent`.
- `boolean nextPosting()`: moves to the next posting (increasing the `index_pointer`) within the current block or reads the next block if necessary. Returns true if a new posting is successfully moved to, false if the end of the posting list is reached.
- `int getTermFrequency(int doc_id)`: retrieves the term frequency associated with the specified document ID in the current block. If the document ID is not found, it returns 0.

- `Integer nextGEQ(int doc_id)`: retrieves the term frequency associated to the next document ID that is greater than or equal to the specified `doc_id`. Returns the term frequency associated to the `doc_id` document, or 0 if not found: note that `getTermFrequency` is used internally.

2.4.3 LexiconReader

It just implements `writeLexicon` and `readLexicon` methods.

2.4.4 DocIndexReader

The `DocIndexReader` class provides utility methods for reading and writing `DocIndex` and `DocInfo` structures from/to disk. Besides the `writeDocIndex` and `readDocIndex` methods, we implemented

- `DocInfo readDocInfo(int doc_id)`: reads `DocInfo` for a specific document ID from the `DocIndex` file.
- `List<Integer> getPids(List<Integer> doc_ids)`: retrieves pids for a list of document IDs.
- `int readN(String doc_index_filename)`: reads the number of entries (N) in a `DocIndex`, but since the `docIndex` file is sorted by `doc_id`, and the `doc_ids` are incrementally, N is the total number of documents.

2.5 SPIMI Algorithm

The SPIMI algorithm is made up by two phases:

- Generate intermediate indexes for each chunk of documents. The chunk size is setted hardcoded to 20.000 documents¹ and each chunk is processed by a thread in order to speed up computation. Then each thread writes its own posting lists, lexicons and document indexes to disk.
- Merging all intermediate indexes generated in the previous phase. First merges intermediate lexicons together, producing for each term a list of `TermEntry` objects: this will speed up computation during the merge of posting lists, since for each term has already the position of its posting lists among all intermediate files. The document index merging is simply the concatenation of the intermediate ones.

¹To avoid `OutOfMemoryError` during execution, the program will wait 1 second if the free memory amount is lower than 20%.

Chapter 3

Experimental Results

All experiments and benchmarks detailed in this chapter were executed on the same machine (MSI Prestige 14 Evo A11M) in order to ensure a standardized and well-defined environment for our experiments. The machine used for benchmarks have the following hardware:

- OS: Microsoft Windows 11 Home 64 bit Ver.2009(OS build 22000.675)
- CPU: 11th Gen Intel(R) Core(TM) i7-1185G7 @ 3.00GHz
- Memory: 16 GB @ 2133 MHz, 8 × 2048 MB, LPDDR4-4267
- Graphics: Intel(R) Iris(R) Xe Graphics, 1024 MB
- Disk: SSD, SAMSUNG MZVL2512HCJQ-00B00, 476,94 GB

Also, all the results are expressed in terms of both mean values and standard deviations, which are derived from 5 independent executions of the experiments. Also, the `README.md` file specifies the commands to compile and execute the program.

Index Construction

The table below summarizes index construction phase, where the flags `-p` (pre-processing) and `-c` (compression) indicate respectively if stemming and stop-word removal must be applied, and if index must be compressed. Also, Index, Lexicon and DocIndex are the size (in MB) of the respective data structures written on disk.

flags	# Terms	Index (MB)	Lexicon (MB)	DocIndex (MB)	Time Elapsed
none	1.369.123	2690	54.8	101	00:15:09 \pm 01:30
-c	1.369.123	1340	54.8	101	00:17:38 \pm 02:22
-p	1.170.498	1430	46.3	101	00:06:30 \pm 01:34
-p -c	1.170.498	738	46.3	101	00:07:35 \pm 01:16

Table 3.1: Index construction performance

Query Execution

The following results are generated issuing the Search Engine on all the 200 queries provided in the `msmarco-test2020-queries.tsv.gz` benchmark collection and only the first `k=20` results are returned.

mode	score	MAP	P@20	NDCG	Time Elapsed (s)
DAAT Conj.	TFIDF	0.139	0.329	0.262	0.023 \pm 0.020
DAAT Conj.	BM25	0.141	0.351	0.266	0.021 \pm 0.017
DAAT Disj.	TFIDF	0.125	0.355	0.245	0.034 \pm 0.030
DAAT Disj.	BM25	0.174	0.448	0.311	0.035 \pm 0.030
MaxScore	TFIDF	0.132	0.367	0.257	0.028 \pm 0.026
MaxScore	BM25	0.182	0.460	0.323	0.026 \pm 0.024

Table 3.2: Index construction performance

3.1 In-Memory Structures

In order to speed up query execution, we decided to keep the whole lexicon in main memory. The below table summarizes the retained size¹ of the lexicon in main memory estimated by IntelliJ Profiler. However to have detailed information, consult the two memory snapshots saved in `data/analysis` folder.

flag	Lexicon (MB)	Total (MB)
none	452.14	487.57
-p	388.43	423.86

Table 3.3: Retained Size of Lexicon if loaded in memory

¹The size of the object plus the size of all objects referenced only by the first object, recursively.